

memory mapped files

- map a file into virtual memory (heap)
- no more `read()` or `write()`, just ordinary memory accesses
- map a file into the address space of a process - the file is mapped into virtual memory
- simplifies file access by treating file I/O through memory rather than `read()/write()` system calls
- page faults may read a page of file data from disk to memory
- allows several processes to map the same file, allowing pages in memory to be shared
 - permits different processes to communicate very efficiently
- requires synchronization between processes that are storing/fetching information to/from the [shared memory](#) region
- faster when copying one file to another

mmap()

- creates a new mapping in the virtual address space of the calling process

```
void * mmap(void *addr, size_t len, int prot, int flags, int fd, off_t offset);
```

- returns the address of the new mapping
- addr
 - NULL: kernel chooses the address at which to create the mapping (most portable)
 - not-NULL: kernel takes it as a hint - Linux will map at a nearby page boundary
- length, offset and fd: initialisation uses `length` bytes starting at `offset` in the file referred to by the file descriptor `fd`
- prot: describes the desired memory protection of the mapping
 - `PROT_NONE` or bitwise OR (`|`) of one or more of the following:
 - `PROT_EXEC` - Pages may be executed
 - `PROT_READ` - Pages may be read
 - `PROT_WRITE` - Pages may be written
 - `PROT_NONE` - Pages may not be accessed
- flags: determines whether updates to the mapping are visible to other processes mapping the same region, and whether updates are carried through to the underlying file
 - MAP_SHARED - Updates to the mapping are visible to other processes that map this file and are carried through to the underlying

file. The file may not actually be updated until `msync` or `munmap` is called.

- `MAP_PRIVATE` - Create a private copy-on-write mapping. Updates to the mapping are not visible to other processes mapping the same file. It is unspecified whether changes made to the file after the `mmap()` call are visible in the mapped region.

- `MAP_ANONYMOUS` - The mapping is not backed by any file; its contents are initialized to zero (e.g. to share a mapped region between two processes).

- `MAP_FIXED` - Does not interpret ``addr`` as a hint: place the mapping at exactly that address. ``addr`` must be a multiple of the page size.

- `MAP_POPULATE` - Populates page tables for a mapping. For a file mapping, this causes read-ahead on the file. Later accesses to the mapping will not be blocked by page faults.

Warning

`offset` must be a multiple of the page size as returned by `sysconf(_SC_PAGE_SIZE)`.

Warning

`prot` must not conflict with the open mode of the file

Note

- Use of a mapped region can result in these signals:
 - `SIGSEGV` - Attempted write into a region mapped as read-only.
 - `SIGBUS` - Attempted access to a portion of the buffer that does not correspond to the file (for example, beyond the end of the file, including the case where another process has truncated the file).

Warning

- Memory mapped by `mmap()` is preserved across `fork()`, with the same attributes.
- A file is mapped in multiples of the page size. For a file that is not a multiple of the page size, the remaining memory is zeroed when mapped, and writes to that region are not written out to the file.
- The effect of changing the size of the underlying file of a mapping, on the pages that correspond to added or removed regions of the file, is unspecified.

- A file cannot be appended with mmap. The file size must be changed first.
- Closing the file descriptor of the mapped file does not unmap the file from memory.

mmap2()

- works like [mmap\(\)](#), but the final argument (`pgoffset`) is expressed in 4096-byte units - enables applications that use 32-bit `off_t` to map large files (up to 2^{44} bytes)

```
void *mmap2(void *addr, size_t length, int prot, int flags, int fd, off_t pgoffset);
```

munmap()

- deletes the mappings for the specified address range

```
int munmap(void *addr, size_t length);
```

- unmap an entire mapping:
 - ``addr``: address returned by ``mmap()``
 - ``length``: the value in the ``mmap()`` call
- unmap part of a mapping:
 - shrinks or cuts it in two, depending on where the unmapping occurs
- regions are automatically unmapped when the process is terminated
- ``addr`` must be a multiple of page size
- all pages containing a part of the indicated range will be unmapped
- subsequent references to these pages will generate SIGSEGV
- it IS NOT an error if the indicated range does not contain any mapped pages

msync()

- flushes the changes made in memory to the underlying file

- updates the part of the file that corresponds to the memory area, starting at `addr` and having a length `length`

```
int msync(void *address, size_t length, int flags);
```

– flags:

– `MS_SYNC` – This flag makes sure the data is actually written to disk. Normally, `msync` only makes sure that accesses to a file with conventional I/O reflect the recent changes.

– `MS_ASYNC` – This tells `msync()` to begin the synchronization, but not to wait for it to complete.

– returns 0 for success, -1 for error

Warning

- in shared mappings, it is the kernel that decides when to write to the underlying file
- without this call, there is no guarantee that changes are written back before [`munmap\(\)`](#) is called

sysconf()

- Memory mapping only works on entire pages of memory.
- Addresses for mapping must be page-aligned, and length values will be rounded up.

To determine the size of a page, use:

```
size_t page_size = (size_t) sysconf(_SC_PAGE_SIZE);
```