

signals

characteristics

- a signal is a software interrupt
- signals notify a process that an event has occurred
- hardware can also create events, which are detected by the HW, received by the kernel and sent to the process as signals (e.g., division by zero)
- **two main categories**
 - *Standard POSIX reliable signals (POSIX.1 standard)*
 - *POSIX real-time signals (POSIX.1b standard)*
 - provide additional signals
 - signals are queued: if multiple instances of a signal are sent, the signal will be received many times
 - data may be sent together with the signals
- each signal is identified by a positive integer, defined in `<signal.h>`
- a signal is **generated by an event** and **delivered to a process**. in between the signal is said **to be pending**
- a pending signal is delivered to a process as soon as it is scheduled to run, or immediately if it is already running

some signals

name	description	default action
SIGALARM	timer expired	terminate
SIGBUS	memory access error	core dumped + terminate
SIGINT	terminal interrupt character (Ctrl + C)	terminate
SIGTSTP	terminal stop character (Ctrl + Z)	stop process
SIGPIPE	program tried to write to a closed channel	
SIGTERM	system requests the application to terminate	
SIGKILL	sure kill signal	terminate

signal delivery

default results

- the signal is ignored - discarded by the kernel and has no effect on the process
- the process is terminated
- a core dump file is generated and the process is terminated
- the process is stopped

what can the process do?

- **ignore the signal**
 - signal is discarded
- **block the signal**
 - signals are stored in a queue (not always) until the process unblocks them, and then they are delivered
- **handle (catch) the signal**
 - signals are redirected to a signal handler which is called
- **none of the above**
 - let the default action apply

Warning

some signals cannot be ignored or handled

`SIGKILL` always terminates the process

`SISTOP` always stops a process

Note

- When a process starts, signals are on their “default behaviour”. Some are ignored, but most are in the “non-handled, non-blocked nor ignored state”. If a signal occurs, the process will die.
- Multiple occurrences of the same signal while a process has blocked it may not result in multiple deliveries when the signal is unblocked
 - POSIX.1 allows the delivery of the signal either once or more than once
 - Only queued signals guarantee the delivery of multiple occurrences
 - Signals are always queued when POSIX real-time extensions are used

basic functions

```
typedef void (*sighandler_t)(int);    // prototype of the handler routine

// receives an integer and returns nothing
sighandler_t signal(int signo, sighandler_t handler);
```

- signo: signal number/name
- handler:
 - can be a function: address to a function that will handle the signal
 - can be `SIG_IGN`: ignores the signal
 - can be `SIG_DFL`: restaures the dafault handler os the signal
- returns the precious handler of the signal specified or SIG_ERR on error
- It is the original API for setting the disposition of a signal and has a simple interface. However, there are variations in the implementation of signal() across different versions of Unix. The function sigaction() does not have that problems and has more functionalities – use it instead of signal(). In some systems signal() is based on sigaction() – in this case it presents no problems

```
int kill(pid_t pid, int signo);
// sends a signal to a certain process
```

- pid: if 0, sends to all processes in the current sender process gorup (including the sender)
- signo: if 0 (`null signal`), no signal is sent and `kill()`only checks if a process can be signaled – can be used to test if a process exists
- returns 0 on success, or -1 on error

```
int raise(int signo);
// send a signal to the process itself
// equivalent to kill(getpid(), signo);
```

- returns 0 on success, or non-zero on error



Note

```
if using pthreads, raise can be implemented as pthread_kill(pthread_self(),  
signo);
```

```
int pause(void);  
// suspends the execution of the process until a signal is received
```

– always returns -1 with `errno` set to `EINTR`

signal sets functions

- a signal set represents multiple signals
- is used in signal functions that need signal sets
- return 0 if successful, -1 on error

```
int sigemptyset(sigset_t *set);  
// initialize the set and exclude all signals  
  
int sigfillset(sigset_t *set);  
// initialize the set and include all signals  
  
int sigaddset(sigset_t *set, int signo);  
// add signal to set  
  
int sigdelset(sigset_t *set, int signo);  
// delete signal from set
```

sigprocmask function

```
int sigprocmask(int how, const sigset_t *restrict set, sigset_t *restrict  
oSet);  
// the signal mask of a process is the set of signals currently blocked  
from delivery to that process
```

```
// sigprocmask() changes the process signal mask (`set` parameter),  
retrieves the existing mask (`oset` parameter), or both
```

– how:

- SIG_BLOCK – signals specified in the signal set pointed to by set are added to the existing signal mask
- SIG_UNBLOCK – signals in the signal set pointed to by set are removed from the existing signal mask
- SIG_SETMASK – the signal set pointed to by `set` replaces the existing signal mask
- returns 0 on success, -1 on error

Info

use `pthread_sigmask()` to manipulate threads signal masks

sigaction function

```
int sigaction(int signo, const struct sigaction *act, struct sigaction  
*oact);  
// sets or examines the action associated to the signal specified
```

- act: if non-NULL, it modifies the action
- oact: if non-NULL, the system returns the previous action for the signal
- returns 0 on success, or -1 on error

```
struct sigaction {  
    void (*sa_handler)(int);    // addr of signal handler (or  
SIG_IGN/SIG_DFL)  
    sigset_t sa_mask;           // signals to block during handler  
    int sa_flags;               // signal options (see manual)  
    void (*sa_restorer)(void); // not for application use  
};
```

- sa_handler: specifies the action to be associated with signal and may be SIG_DFL for the default action, SIG_IGN to ignore this signal, or a

pointer to a signal handling function

- `sa_mask`:

- defines the signals to be blocked during the invocation of the handler `sa_handler`

- the signal that caused the handler to be invoked is automatically added to the mask (is also blocked)

- when the handler returns the `sa_mask` is removed

- it specifies the set of signals that cannot interrupt this handler

other functions

- other functions
 - `sigismember()` : test for membership of a set.
 - `sigandset()` : intersects sets (and)
 - `sigorset()` : union of sets (or)
 - `sigpending()` : to determine which signals are pending for a process
 - `killpg()` : send signal to a process group
 - ...
- for threads
 - `pthread_sigmask()` : signal mask for threads
 - `sigwait()` : wait for a signal
 - `pthread_kill()` : send a signal to a thread

some considerations

- after calling a standard function, it may return -1 indicating an error
 - `errno == EINTR` means that a certain routine was interrupted and has to be tried again.
 - other routines return other things.
 - if you are using signals, you must protect them against all that!
- What do you think it will happen if you receive a signal inside a signal handler??
 - in most systems, upon entering a signal handling routine, all signals of that type become blocked (i.e. they are queued - for “normal” signals, a finite set of them are queued (typically 1); for “real time signals”, all are...)

- the other signals are still processed asynchronously if they arrive.
- this behaviour is not consistent across systems. In fact, in some systems, that signal type resets to its default behaviour. This means that if, meanwhile, the program receives a signal of the same type it may die! On that type of system, the first thing that you must do is to once again set the signal handler.
- the new POSIX routines address this. Also, most system nowadays do not reset the signal handler.

Signal	Architecture			Standard that specified the signal	Action	Comment
	x86/ARM most others	Alpha/ SPARC	MIPS			
SIGABRT	6	6	6	P1990	Core	Abort signal from abort(3)
SIGALRM	14	14	14	P1990	Term	Timer signal from alarm(2)
SIGBUS	7	10	10	P2001	Core	Bus error (bad memory access)
SIGCHLD	17	20	18	P1990	Ign	Child stopped or terminated
SIGCLD	-	-	18	-	Ign	A synonym for SIGCHLD
SIGCONT	18	19	25	P1990	Cont	Continue if stopped
SIGEMT	-	7	7	-	Term	Emulator trap
SIGFPE	8	8	8	P1990	Core	Floating-point exception
SIGHUP	1	1	1	P1990	Term	Hangup detected on controlling terminal or death of controlling process
SIGILL	4	4	4	P1990	Core	Illegal Instruction
SIGINFO	-	29/-	-	-	-	A synonym for SIGPWR
SIGINT	2	2	2	P1990	Term	Interrupt from keyboard (^C)
SIGIO	29	23	22	-	Term	I/O now possible (4.2BSD)
SIGIOT	6	6	6	-	Core	IOT trap. A synonym for SIGABRT
SIGKILL	9	9	9	P1990	Term	Kill signal (cannot be caught, blocked or ignored)
SIGLOST	-	-/29	-	-	Term	File lock lost (unused)
SIGPIPE	13	13	13	P1990	Term	Broken pipe: write to pipe with no readers; see pipe(7)
SIGPOLL	-	-	-	P2001	Term	Pollable event (Sys V); synonym for SIGIO
SIGPROF	27	27	29	P2001	Term	Profiling timer expired
SIGPWR	30	29/-	19	-	Term	Power failure (System V)
SIGQUIT	3	3	3	P1990	Core	Quit from keyboard
SIGSEGV	11	11	11	P1990	Core	Invalid memory reference (segmentation violation)
SIGSTKFLT	16	-	-	-	Term	Stack fault on coprocessor (unused)
SIGSTOP	19	17	23	P1990	Stop	Stop process (cannot be caught, blocked or ignored)
SIGSYS	31	12	12	P2001	Core	Bad system call (SVr4); see also seccomp(2)
SIGTERM	15	15	15	P1990	Term	Termination signal
SIGTRAP	5	5	5	P2001	Core	Trace/breakpoint trap
SIGTSTP	20	18	24	P1990	Stop	Stop typed at terminal (^Z)
SIGTTIN	21	21	26	P1990	Stop	Terminal input for background process
SIGTTOU	22	22	27	P1990	Stop	Terminal output for background process
SIGUNUSED	31	-	-	-	Core	Synonymous with SIGSYS
SIGURG	23	16	21	P2001	Ign	Urgent condition on socket (4.2BSD)
SIGUSR1	10	30	16	P1990	Term	User-defined signal 1
SIGUSR2	12	31	17	P1990	Term	User-defined signal 2
SIGVTALRM	26	26	28	P2001	Term	Virtual alarm clock (4.2BSD)
SIGWINCH	28	28	20	-	Ign	Window resize signal (4.3BSD, Sun)
SIGXCPU	24	24	30	P2001	Core	CPU time limit exceeded (4.2BSD); see setrlimit(2)
SIGXFSZ	25	25	31	P2001	Core	File size limit exceeded (4.2BSD); see setrlimit(2)

P1990 = signal described in POSIX.1-1990 standard

• signals and [threads](#)

- Some conflicts exist between the Unix signal model that was based on processes and POSIX threads model.
- Signal actions are process-wide – if a signal received by a thread has as default action to stop or terminate, then all threads in the process are stopped or terminated.

- Actions are shared between all threads – when an action is changed in a thread, all threads response to the signal is changed.
- A signal may be directed to a specific thread (thread-directed):
 - in the case of synchronous signals that result from a specific thread execution, which are received by that same thread (e.g., SIGFPE – floating point exception);
 - by using `pthread_kill()` to enable a thread to send a signal to other thread;
- A signal is process-directed:
 - if sent from a process to other process;
 - in the case of asynchronous signals, which are received by any thread that has not blocked them (`pthread_sigmask`);
- When a signal is delivered to a multithreaded process just one thread catches it
- Each thread may have its own thread mask
- Signals from the Linux shell
 - Using Linux command line
 - List signals: `kill -l`
 - Send a `SIGALRM` to a process with `PID = 76543`: `kill -SIGALRM 76543`
 - Send a signal to all processes with the name `myproc`: `killall -SIGKILL myproc`
 - CTRL + C
 - When '^C' is pressed the terminal will send a SIGINT to the foreground process group of the terminal